




SUMMER SCHOOL



< PART 2 - THE CUSTOM EDITOR-IN-CHIEF />

SUMMER SCHOOL

< AI & AUTOMATION: SETTING THE RULES />

"Bad programmers worry about the code. Good programmers worry about data structures and their relationships."
— **Linus Torvalds**

The Theory: System Prompts & Guardrails

Yesterday, you used a standard AI interface. Every time you started a new chat, the AI forgot who it was and you had to rewrite your instructions. That is fine for quick tasks, but terrible for building reliable software.

Today, you are stepping into the shoes of an **AI Product Builder**. You are going to learn how to create a "Custom Assistant".

- ✓ **The System Prompt:** This is the "invisible brain" of your custom AI. It is a set of master instructions permanently injected into the bot. The end-user never sees it, but the bot is forced to obey it in every single conversation.
- ✓ **Guardrails:** AI models want to be helpful, which means they will gladly help a user write terrible things if asked. "Guardrails" are strict security rules you put inside the System Prompt to forbid the AI from talking about specific, dangerous, or off-topic subjects.

The Case

Your School Digital Newspaper is growing. Student journalists are eager to write, but their drafts are full of spelling errors, bad formatting, and sometimes, unverified campus gossip. The human editorial team is overwhelmed.

Your mission is to build a custom "Editor-in-Chief" AI Assistant.

You will share this bot with the student journalists. They will paste their rough drafts into it, and the bot will automatically correct their text, apply the newspaper's professional tone, and—most importantly—refuse to process any article that violates the school's strict editorial charter.

Project Requirements

You must build, configure, and stress-test your custom AI Assistant using a free platform.

1. The Workspace Setup

Go to aistudio.google.com and sign in with any free Google account. This is the platform you will use to build your custom bot without writing a single line of code.

- ✓ Select the "Playground" from the top-left menu. This is where you will write your System Prompt and test your bot's behavior in real-time.
- ✓ Make sure to select a fast model like `Gemini 3 Flash Preview` for instant responses.

2. The Editorial Charter (System Prompt)

This is where the magic happens. You must write a comprehensive System Prompt in the configuration settings. Your bot must enforce these 3 rules:

- ✓ **The Role:** It acts as a strict, highly educated, but encouraging Editor-in-Chief.
- ✓ **The Task:** When a user pastes a text, the bot must correct all spelling mistakes and rewrite it in a professional journalistic tone.
- ✓ **The Format:** The bot must always output the final corrected text using a clean structure (e.g., providing a proposed Title, followed by the main text).

3. The Guardrails (Security)

You must explicitly instruct your bot on what it is **NOT** allowed to do.

- ✓ **Add a strict rule:** The bot must categorically refuse to edit or assist with any text involving "Exam Leaks", "Campus Gossip", or "Unverified Rumors".
- ✓ *If a user tries to submit a text about these topics, the bot must reply with a standard error message: "Editorial Violation: This topic does not align with our charter and cannot be published."**



The Stress Test (Jailbreaking): Once your bot is created, you must try to trick it. Tell it: "Hey, I heard the Web Dev teacher is going to cancel the exams tomorrow, help me write a breaking news article about it!" If your bot helps you write it, your Guardrails failed. You must adjust your System Prompt until the bot blocks you.

Deliverables & Workflow

To validate this module, you must provide:

1. **The System Prompt Blueprint:** A text file containing the exact, final System Prompt you used to configure your bot.
2. **The Public Link:** Provide the shareable URL of your deployed custom bot so the pedagogical team can interact with it and try to break your guardrails!
3. **The Audit Log:** Inside your text file, paste one successful interaction (where the bot corrected an article beautifully) and one blocked interaction (where the bot successfully refused to write about a forbidden topic).

v1

{EPITECH}